

hw6-25912237

November 23, 2015

All relevant code is in this notebook and two files `neuralnet.py` and `loader.py`.

```
In [50]: %matplotlib inline
         %load_ext autoreload
         import numpy as np
         import pandas as pd
         import scipy as sp
         import scipy.io
         import datetime
         from sklearn.cross_validation import train_test_split
         from sklearn.preprocessing import normalize
         from sklearn.utils import shuffle

         import matplotlib.pyplot as plt
         import matplotlib
         %matplotlib inline
```

The autoreload extension is already loaded. To reload it, use:
`%reload_ext autoreload`

```
In [51]: import loader
         import neuralnet
         %aimport neuralnet
         %aimport loader
         %aimport
```

Modules to reload:
loader neuralnet

Modules to skip:

1 Problem 1

\$ y_0 = 784 \times 1, y_1 = 200 \times 1, y_2 = 10 \times 1 \$

$$\begin{aligned} z_1 &= 200x_1, z_2 = 10x_1 \\ b_1 &= 200x_1, b_2 = 10x_1 \\ w_1 &= 784x_{200}, w_2 = 200x_1 \\ g &= \frac{1}{1 + e^{-z}} \\ t &= \tanh(z) \end{aligned}$$

Rather than adding a neuron, I've elected to just add an arbitrary, updating value to each z . Which is equivalent to adding another neuron.

$$\frac{\partial L}{\partial w^l} = y^{l-1} \frac{\partial L}{\partial z^l}^T$$

$$\begin{aligned}
\frac{\partial L}{w^2} &= y^1 \frac{\partial L}{\partial z^2}^T \\
&= y^1 (\text{diag}(g'(z^2 + b^2))) \cdot \frac{\partial L}{\partial y^2}^T \\
w^2 &\leftarrow w^2 - \eta * \frac{\partial L}{w^2} \\
\frac{\partial L}{w^1} &= y^0 \frac{\partial L}{\partial z^1}^T \\
&= y^0 (\text{diag}(t'(z^1 + b^1))) \cdot \frac{\partial L}{\partial y^1}^T \\
&= y^0 (\text{diag}(t'(z^1 + b^1))) \cdot w^2 \cdot \frac{\partial L}{\partial z^2}^T \\
w^1 &\leftarrow w^1 - \eta * \frac{\partial L}{w^1}
\end{aligned}$$

The benefit of backprop is being able to cache the results from the previous delta.

1.0.1 Squared Loss:

$$\begin{aligned}
\frac{\partial L}{w^l} &= y^0 (\text{diag}(1 - \tanh^2(z^1 + b^1))) \cdot w^2 \cdot (\text{diag}(g(z^2 + b^2)(1 - g(z^2 + b^2)))) \cdot (y^2 - y))^T \\
\frac{\partial L}{w^2} &= y^1 (\text{diag}(g(z^2 + b^2)(1 - g(z^2 + b^2)))) \cdot (y^2 - y))^T
\end{aligned}$$

1.0.2 Cross-Entropy Loss:

$$\begin{aligned}
\frac{\partial L}{w^l} &= y^0 (\text{diag}(1 - \tanh^2(z^1 + b^1))) \cdot w^2 \cdot (\text{diag}(g(z^2 + b^2)(1 - g(z^2 + b^2)))) \cdot -(\frac{y}{y^2} - \frac{1 - y}{1 - y^2}))^T \\
\frac{\partial L}{w^2} &= y^1 (\text{diag}(g(z^2 + b^2)(1 - g(z^2 + b^2)))) \cdot -(\frac{y}{y^2} - \frac{1 - y}{1 - y^2}))^T
\end{aligned}$$

```

In [52]: def squared_error(y, y_hat):
          return (y - y_hat).dot((y - y_hat))

def squared_error_prime(y, y_hat):
    return -1 * (y - y_hat)

def cross_ent(y, y_hat):
    part1 = y.dot(np.log(y_hat))
    part2 = (1.0 - y).dot(np.nan_to_num(np.log(1.0 - y_hat)))
    return -(part1 + part2)

def cross_ent_prime(y, y_hat):
    part1 = y / y_hat
    part2 = (1 - y)/(1 - y_hat)
    return -(part1 - part2)

def t(z):
    return np.tanh(z)

def t_prime(z):

```

```

        return 1 - np.power(t(z), 2)

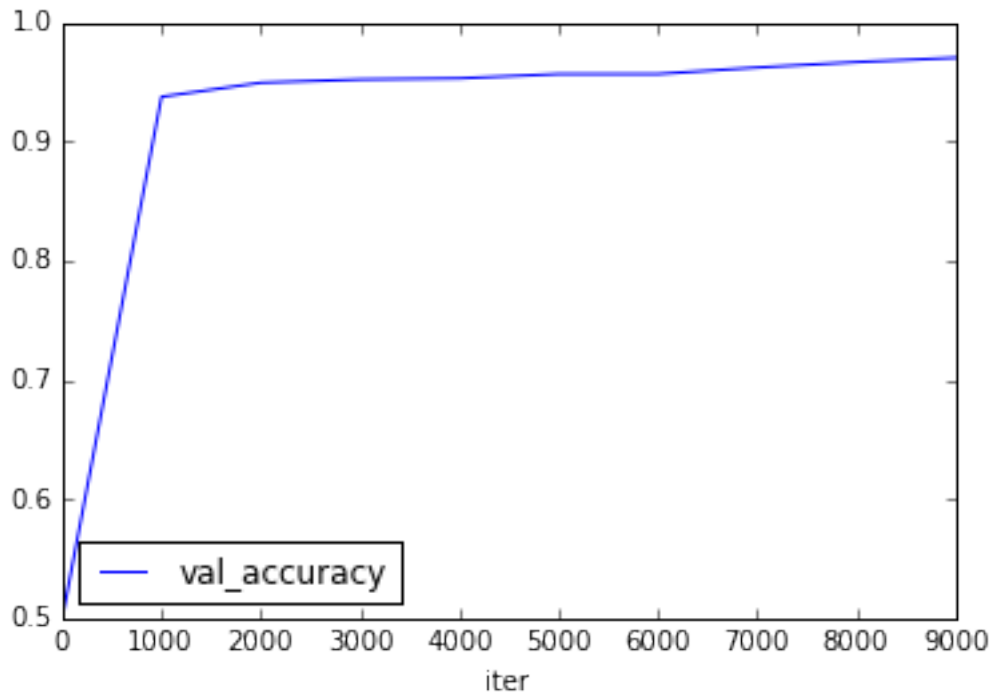
def g(z):
    return (1.0 / (1.0 + np.exp(-z)))

def g_prime(z):
    return g(z) * (1 - g(z))
In [55]: %autoreload
        input_dim = len(X[0])
        output_dim = len(y[0])
        hl_size = 50
        layer1 = neuralnet.Layer((input_dim, hl_size), t, t_prime)
        layer2 = neuralnet.OutputLayer((hl_size, output_dim), g, g_prime)
        nn = neuralnet.NeuralNet(layer1, layer2, cross_ent, cross_ent_prime)
        nn.train(X, y, XVal=X, yVal=y, num_iters=10000)
        nn.resulting_scores()[['val_accuracy', 'iter']].plot(x='iter')
        print("Parameters used", nn.train_params)
        print("Training Accuracy", nn.score(X,y)/len(X))

Total Train Time 1.48 Seconds
Parameters used {'end': datetime.datetime(2015, 11, 23, 13, 41, 26, 357580), 'start': datetime.datetime(2015, 11, 23, 13, 41, 26, 357580)}
Training Accuracy 0.969090909091

/Users/bill_chambers/AeroFS/Dev/ml289/hw6/neuralnet.py:109: FutureWarning: comparison to 'None' will result in an element-wise comparison in the future.
if XVal != None:

```



2 Problem 2

1. Parameters are inline below, weights are initialized normally around $\mu = 1, \sigma = 0.01$

2. Training and Validation Accuracies are listed below.
3. Train time is listed below.
4. Plots are below. I only report every 20,000 iterations.
5. The cross entropy seemed to perform much better. Simply in terms of the equation it seems to be a more sophisticated measure of the loss. Therefore we could expect it to perform better. This aligned with the results, as it allowed me to descend faster and get to a much smaller error rate overall.

Kaggle Score for 200k iterations, 0.01 learning rate, squared loss = .929

Kaggle Score for 500k iterations, 0.01 learning rate, cross entropy loss = 0.97460

it's interesting to note that these had the exact same initialization values (for w's) Even with 200k, the cross entropy loss was far superior.

```
In [56]: train_dict = loader.load_training_data() # loader.py
        X = train_dict['X']
        y = train_dict['y']
        print("Mean of y:", y.mean())
        y = pd.get_dummies(y).values

        XVal = train_dict['XVal']
        yVal = pd.get_dummies(train_dict['yVal']).values

        print("Num Training Examples:", len(X))
```

Mean of y: 4.44715555556

Num Training Examples: 45000

2.1 Squared Error Loss

```
In [59]: %autoreload
        input_dim = len(X[0])
        output_dim = len(y[0])
        hl_size = 200

        layer1 = neuralnet.Layer((input_dim, hl_size), t, t_prime)
        layer2 = neuralnet.OutputLayer((hl_size, output_dim), g, g_prime)

        nn = neuralnet.NeuralNet(layer1, layer2, squared_error, squared_error_prime)
        nn.train(X, y, num_iters=200000, score_every=20000, eta=0.01, XVal=XVal, yVal=yVal)
```

Total Train Time 352.62 Seconds


/Users/bill_chambers/AeroFS/Dev/ml289/hw6/neuralnet.py:109: FutureWarning: comparison to 'None' will result in an element-wise object comparison if XVal != None:

```
In [60]: print("Parameters used", nn.train_params)
        nn.resulting_scores()[['val_accuracy', 'iter']].plot(x='iter')
        print("Training Accuracy", nn.score(X,y)/len(X))
        print("Validation Accuracy", nn.score(XVal,yVal)/len(XVal))
```

Parameters used {'score_every': 20000, 'start': datetime.datetime(2015, 11, 23, 13, 42, 11, 482967), 'et

Training Accuracy 0.9338

Validation Accuracy 0.9278



2.2 Cross Entropy Loss

```
In [64]: %autoreload
         input_dim = len(X[0])
         output_dim = len(y[0])
         hl_size = 200

         layer1 = neuralnet.Layer((input_dim, hl_size), t, t_prime)
         layer2 = neuralnet.OutputLayer((hl_size, output_dim), g, g_prime)

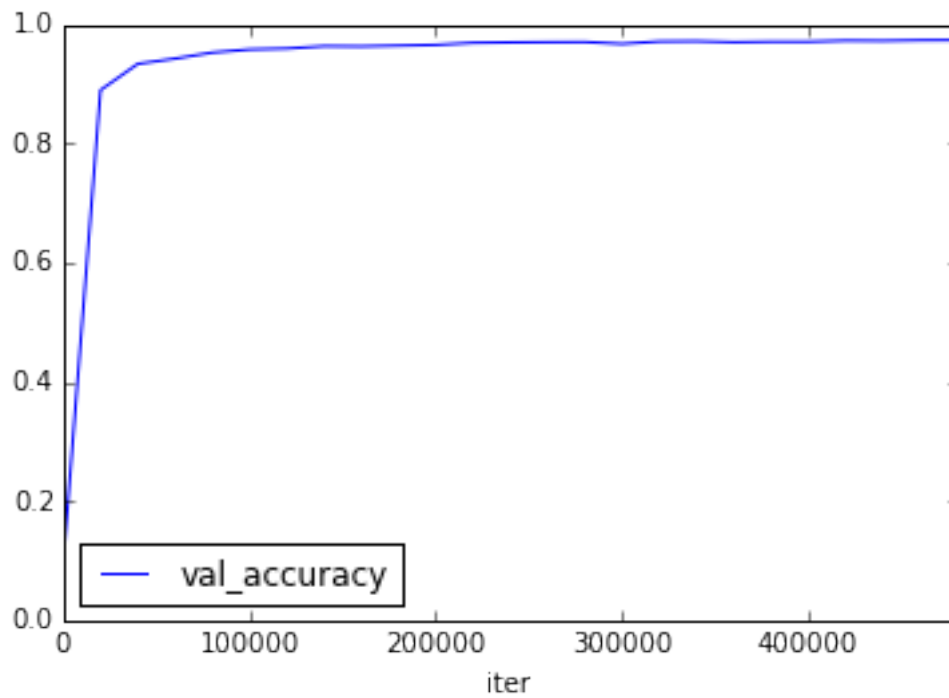
         nn = neuralnet.NeuralNet(layer1, layer2, cross_ent, cross_ent_prime)
         nn.train(X, y, num_iters=500000, score_every=20000, eta=0.01, XVal=XVal, yVal=yVal)
```

Total Train Time 656.30 Seconds

/Users/bill_chambers/AeroFS/Dev/ml289/hw6/neuralnet.py:109: FutureWarning: comparison to 'None' will result in an element-wise object comparison in the future
if XVal != None:

```
In [65]: print("Parameters used", nn.train_params)
         nn.resulting_scores()[['val_accuracy', 'iter']].plot(x='iter')
         print("Training Accuracy", nn.score(X,y)/len(X))
         print("Validation Accuracy", nn.score(XVal,yVal)/len(XVal))
```

Parameters used {'score_every': 20000, 'start': datetime.datetime(2015, 11, 23, 14, 0, 31, 23813), 'eta': 0.01, 'num_iters': 500000, 'XVal': None, 'yVal': None, 'cross_ent': 0.0001, 'cross_ent_prime': 0.0001, 't': 0.0, 't_prime': 0.0, 'g': 0.0, 'g_prime': 0.0, 'hl_size': 200, 'input_dim': 1000, 'output_dim': 1000, 'layer1': neuralnet.Layer((1000, 200), 0.0, 0.0), 'layer2': neuralnet.OutputLayer((200, 1000), 0.0, 0.0)}
Training Accuracy 0.997533333333
Validation Accuracy 0.972333333333



```
In [66]: kaggleTest = loader.load_test_data()
         preds = nn.predict(kaggleTest)
         df = pd.DataFrame(preds).reset_index()
         df.columns = ['id', 'category']
         df.id = df.id + 1
         df.to_csv("cross_ent 500000 iters.csv", index=None)
```