


# Dimensionality Reduction by PCA

# Pattern Matrix




# Examples: Pattern Matrices

$$\mathbf{X} \in \mathbb{R}^{n \times m}$$


## ► Measurement vectors

- $i$ : instance number, e.g. a house
- $j$ : measurement, e.g. the area of a house




## ► Digital images as gray-scale vectors

- $i$ : image number
- $j$ : pixel value at location  $j=(k,l)$




## ► Text documents in bag-of-words representation

- $i$ : document number
- $j$ : term (word or phrase) in a vocabulary



## ► User rating data

- $i$ : user number
- $j$ : item (book, movie)



# Document-Term Matrix

$D$  = Document collection

$W$  = Lexicon/Vocabulary

intelligence

$w_j$

*Texas Instruments said it has developed the first 32-bit computer chip designed specifically for artificial intelligence applications [...]*

$d_i$  = 

...	0	1	...	2	0	...
-----	---	---	-----	---	---	-----

  
X  
term weighting

$t$

term weighting


Document-Term Matrix

	$W$					
	$w_1$	...	$w_j$	...	$w_J$	
$d_1$						
...			...			
$d_i$		...	$c(d_i, w_j)$	...		
...			...			
$d_I$						

# A 100 Million<sup>ths</sup> of a Typical Document-term Matrix

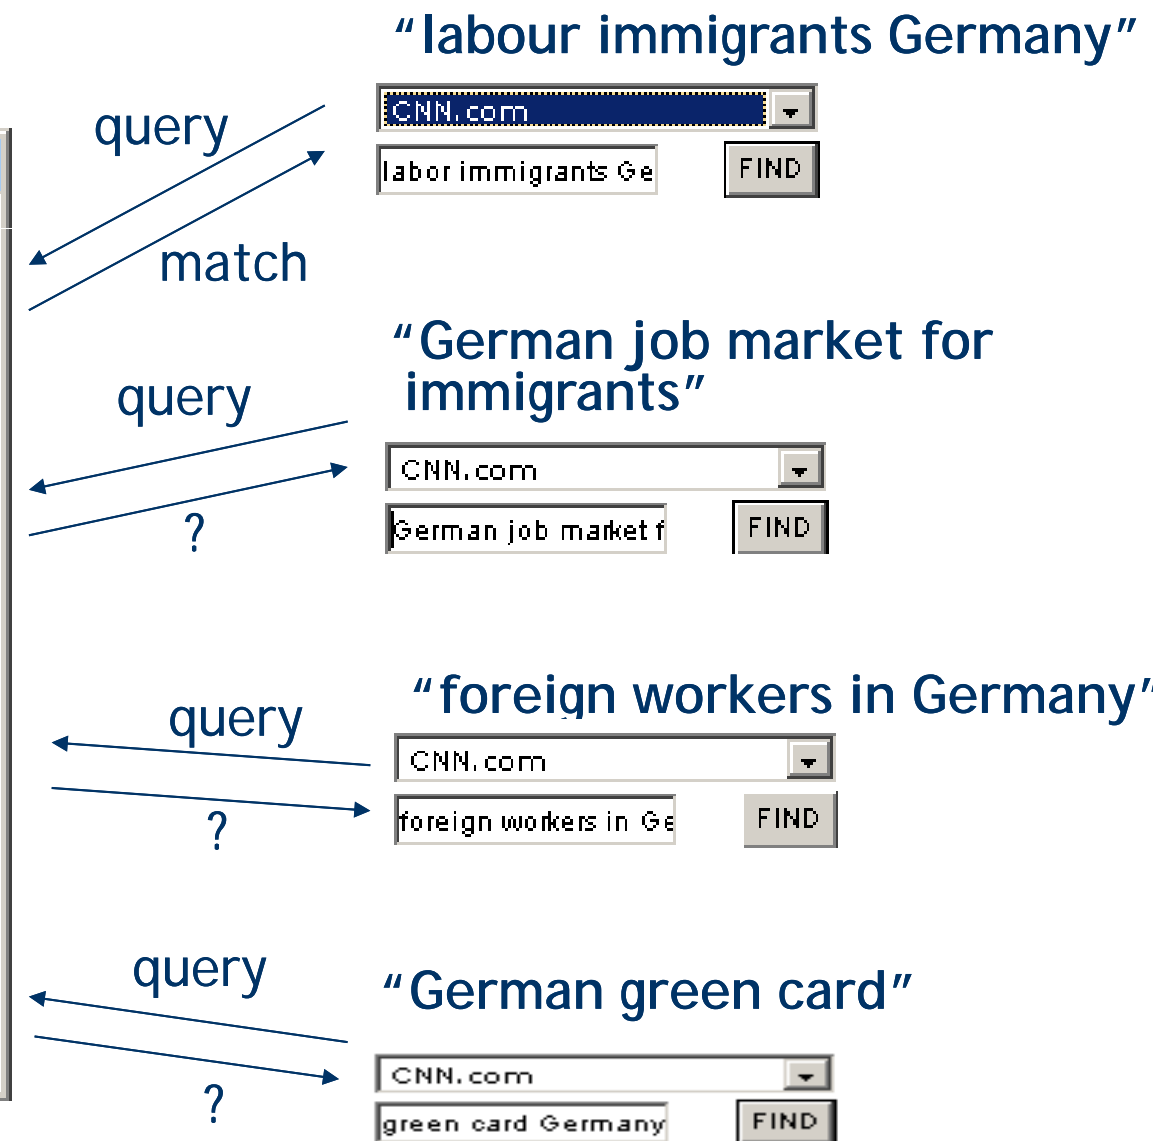
Typical:

- Number of documents  $\approx 1.000.000$
- Vocabulary  $\approx 100.000$
- Sparseness  $< 0.1 \%$
- Fraction depicted  $\approx 1e-8$





# Vocabulary Mismatch & ~~Robustness~~



# Document-Term Matrix

$D$  = Document collection

$W$  = Lexicon/Vocabulary

intelligence

$w_j$

*Texas Instruments said it has developed the first 32-bit computer chip designed specifically for artificial intelligence applications [...]*

$d_i$  = 

...	0	1	...	2	0	...
-----	---	---	-----	---	---	-----

  
X  
term weighting

$t$

term weighting


Document-Term Matrix

	$W$					
	$w_1$	...	$w_j$	...	$w_J$	
$d_1$						
...			...			
$d_i$		...	$c(d_i, w_j)$	...		
...			...			
$d_I$						

Clustering rows?

Clustering columns?

# Latent Structure



- ▶ Given a matrix that “encodes” data ...

- ▶ Potential problems

- too large
- too complicated
- missing entries
- noisy entries
- lack of structure
- ...

$$\mathbf{A} = \begin{pmatrix} a_{11} & \dots & a_{1j} & \dots & a_{1m} \\ \dots & \dots & \dots & \dots & \dots \\ a_{i1} & \dots & a_{ij} & \dots & a_{im} \\ \dots & \dots & \dots & \dots & \dots \\ a_{n1} & \dots & a_{nj} & \dots & a_{nm} \end{pmatrix}$$

- ▶ Is there a **simpler** way to **explain** entries?
- ▶ There might be a **latent structure** underlying the data.
- ▶ How can we “find” or “reveal” this structure?



# Matrix Decomposition

- Common approach: approximately **factorize** matrix

$$\mathbf{A} \approx \hat{\mathbf{A}} = \mathbf{L} \cdot \mathbf{R}$$

approximation      left factor      right factor


- Factors are typically constrained to be “thin”

$$\begin{array}{c} \overbrace{\hspace{2cm}}^m \\ \begin{array}{|c|} \hline \mathbf{A} \\ \hline \end{array} \\ \underbrace{\hspace{2cm}}_n \end{array} \approx \begin{array}{c} \overbrace{\hspace{2cm}}^q \\ \begin{array}{|c|} \hline \mathbf{L} \\ \hline \end{array} \\ \underbrace{\hspace{2cm}}_n \end{array} \cdot \begin{array}{c} \overbrace{\hspace{2cm}}^m \\ \begin{array}{|c|} \hline \mathbf{R} \\ \hline \end{array} \\ \underbrace{\hspace{2cm}}_q \end{array}$$

reduction  
 $n \cdot m \gg n \cdot q + m \cdot q$   
 factors = latent structure (?)


# Dimensionality reduction

- Ex: data with two real values  $[x_1, x_2]$
- We'd like to describe each point using only one value  $[z_1]$
- We'll communicate a "model" to convert:  $[x_1, x_2] \sim f(z_1)$





# Dimensionality reduction

- Ex: data with two real values  $[x_1, x_2]$
- We'd like to describe each point using only one value  $[z_1]$
- We'll communicate a "model" to convert:  $[x_1, x_2] \sim f(z_1)$
- Ex: linear function  $f(z)$ :  $[x_1, x_2] = z * \underline{v} = z * [v_1, v_2]$
- $\underline{v}$  is the same for all data points (communicate once)
- $z$  tells us the closest point on  $v$  to the original point  $[x_1, x_2]$



# Dimensionality reduction

- Ex: data with two real values  $[x_1, x_2]$
- We'd like to describe each point using only one value  $[z_1]$
- We'll communicate a "model" to convert:  $[x_1, x_2] \sim f(z_1)$
- Ex: linear function  $f(z)$ :  $[x_1, x_2] = z * \underline{v} = z * [v_1, v_2]$
- $\underline{v}$  is the same for all data points (communicate once)
- $z$  tells us the closest point on  $v$  to the original point  $[x_1, x_2]$




# Principal Components Analysis



- What is the vector that would most closely reconstruct  $X$ ?

$$\min_{a,v} \sum_i (x^{(i)} - a^{(i)}v)^2$$

- Given  $v$ :  $a^{(i)}$  is the projection of each point  $x^{(i)}$  onto  $v$
- $v$  chosen to minimize the residual variance
- Equivalently,  $v$  is the direction of maximum variance
- Extensions: best two dimensions:  $x_i = a_i v + b_i w + m$



# Eigenvectors



# PCA

Given pattern matrix  $X$ ,

1. Subtract mean from each point
2. (sometimes) scale each dimension by its variance
3. Compute covariance matrix  $C = X^T X$
4. Compute  $k$  largest eigenvectors of  $C$

$$C = V D V^T$$

# Singular Value Decomposition

- Alternative method to calculate (still subtract mean 1<sup>st</sup>)
- Decompose  $X = U S V^T$ 
  - Orthogonal:  $X^T X = V S S V^T = V D V^T$
  - $X X^T = U S S U^T = U D U^T$
- $U \cdot S$  matrix provides coefficients
  - Example  $x_i = U_{i,1} S_{11} V_1 + U_{i,2} S_{22} V_2 + \dots$
- Gives the least-squares approximation to  $X$  of this form

$$\boxed{\begin{matrix} X \\ N \times D \end{matrix}} \approx \boxed{\begin{matrix} U \\ N \times K \end{matrix}} \boxed{\begin{matrix} S \\ K \times K \end{matrix}} \boxed{\begin{matrix} V^T \\ K \times D \end{matrix}}$$



# Glorious SVD

$$X = USV^T$$

- $XX^T$  and  $X^T X$  share the same eigenvalues
- Even better: their eigenvectors are related
  - $Xv_i$  is an eigenvector of  $XX^T$

# Collaborative Filtering (Netflix)

From Y. Koren  
of BellKor team

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1		3		?	5			5		4	
	2			5	4			4			2	1	3
	3	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	6	1		3		3			2			4	

$$\begin{matrix} X \\ N \times D \end{matrix} \approx \begin{matrix} U \\ N \times K \end{matrix} \begin{matrix} S \\ K \times K \end{matrix} \begin{matrix} V^T \\ K \times D \end{matrix}$$

From Y. Koren  
of BellKor team

# Latent Space Models

Model ratings matrix as  
“user” and “movie”  
positions

Infer values from known  
ratings

		users											
items	1		3			5			5		4		
			5	4			4			2	1	3	
	2	4		1	2		3		4	3	5		
		2	4		5			4			2		
			4	3	4	2					2	5	
	1		3		3			2			4		

~

Extrapolate to unranked

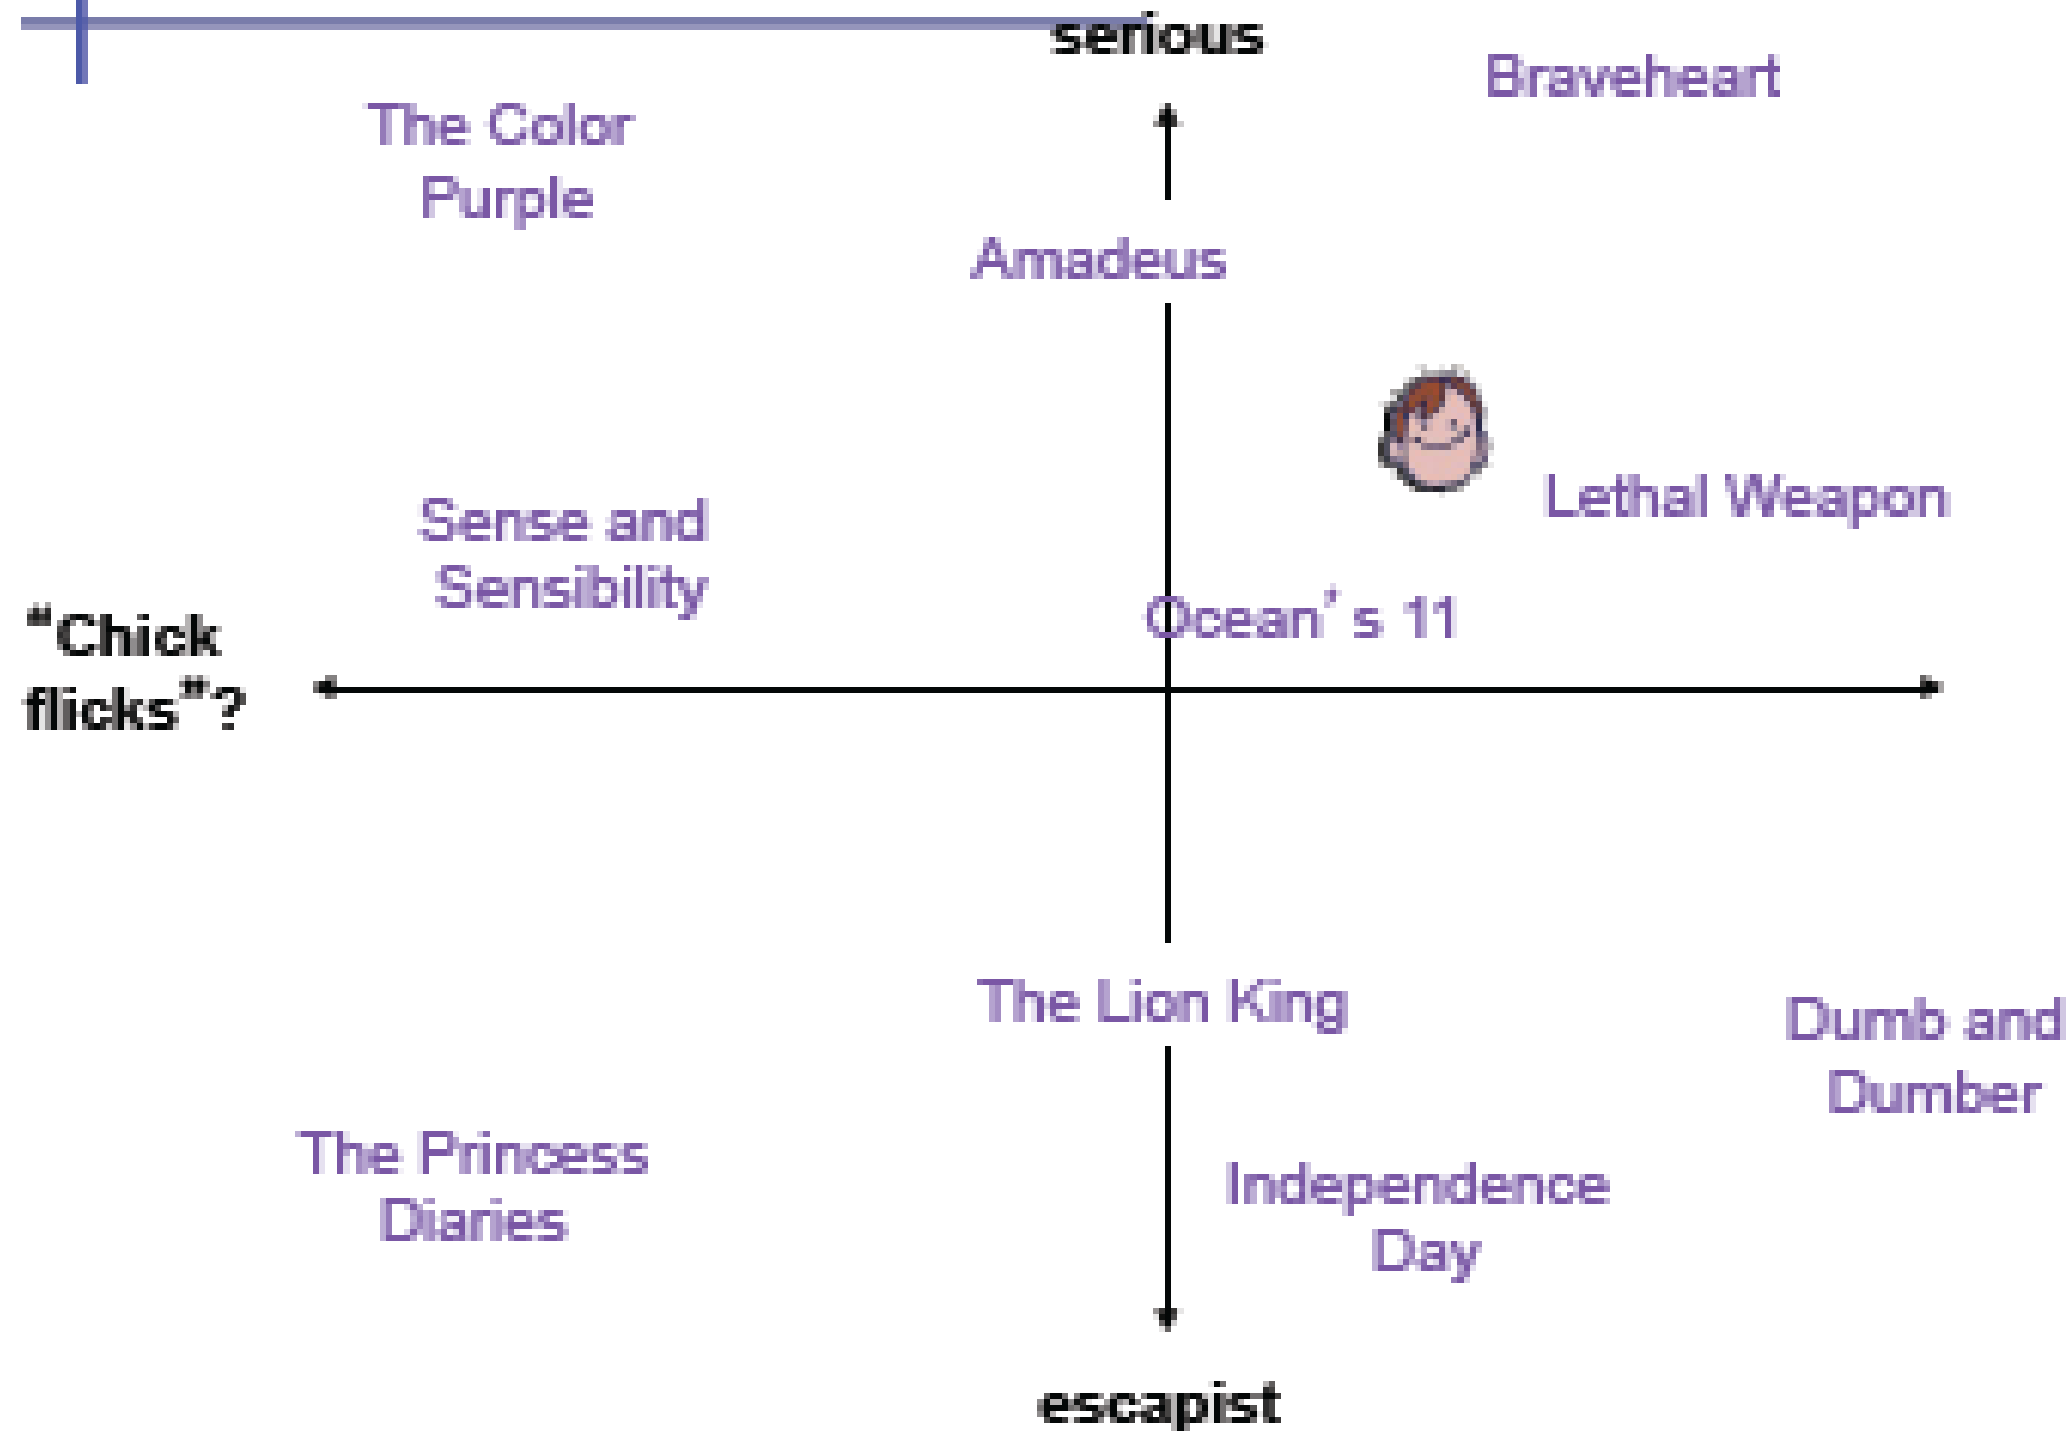
	users										
items	.1	-.4	.2								
	-.5	.6	.5								
	-.2	.3	.5								
	1.1	2.1	.3								
	-.7	2.1	-.2								
	-.1	.7	.3								

•

1.1	-.2	.3	.5	-.2	-.5	.8	-.4	.3	1.4	2.4	-.9
-.8	.7	.5	1.4	.3	-.1	1.4	2.9	-.7	1.2	-.1	1.3
2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1


From Y. Koren  
of BellKor team

# Latent Space Models



# “Eigen-faces”

- “Eigen-X” = represent X using PCA
- Ex: Viola Jones data set
  - 24x24 images of faces = 576 dimensional measurements



# “Eigen-faces”

- “Eigen-X” = represent X using PCA
- Ex: Viola Jones data set
  - 24x24 images of faces = 576 dimensional measurements
  - Take first K PCA components



+



⋮


$$\begin{matrix} X \\ N \times D \end{matrix}$$

$\approx$


$$\begin{matrix} U \\ N \times K \end{matrix}$$

$$\begin{matrix} S \\ K \times K \end{matrix}$$


$$\begin{matrix} V^T \\ K \times D \end{matrix}$$




Mean




V(1,:)



V(2,:)



V(3,:)



V(4,:)


...

# “Eigen-faces”


- “Eigen-X” = represent X using PCA
- Ex: Viola Jones data set
  - 24x24 images of faces = 576 dimensional measurements
  - Take first K PCA components




Mean




Dir 1



Dir 2





Dir 3




Dir 4



...




K=4



K=50




K=4



K=50

# The Face Subspace

---



The set of faces is a “subspace” of the set of images

- Suppose it is  $K$  dimensional
- We can find the best subspace using PCA
- This is like fitting a “hyper-plane” to the set of faces
  - spanned by vectors  $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_K$

$$\text{Any face: } \mathbf{x} \approx \bar{\mathbf{x}} + a_1 \mathbf{v}_1 + a_2 \mathbf{v}_2 + \dots + a_k \mathbf{v}_k$$




# Projecting onto the Eigenface Subspace

---

- The eigenfaces  $\mathbf{v}_1, \dots, \mathbf{v}_K$  span the space of faces
  - A face is converted to eigenface coordinates by


$$\mathbf{x} \rightarrow (\underbrace{(\mathbf{x} - \bar{\mathbf{x}}) \cdot \mathbf{v}_1}_{a_1}, \underbrace{(\mathbf{x} - \bar{\mathbf{x}}) \cdot \mathbf{v}_2}_{a_2}, \dots, \underbrace{(\mathbf{x} - \bar{\mathbf{x}}) \cdot \mathbf{v}_K}_{a_K})$$

$$\mathbf{x} \approx \bar{\mathbf{x}} + a_1 \mathbf{v}_1 + a_2 \mathbf{v}_2 + \dots + a_K \mathbf{v}_K$$



# Choosing the Dimension K

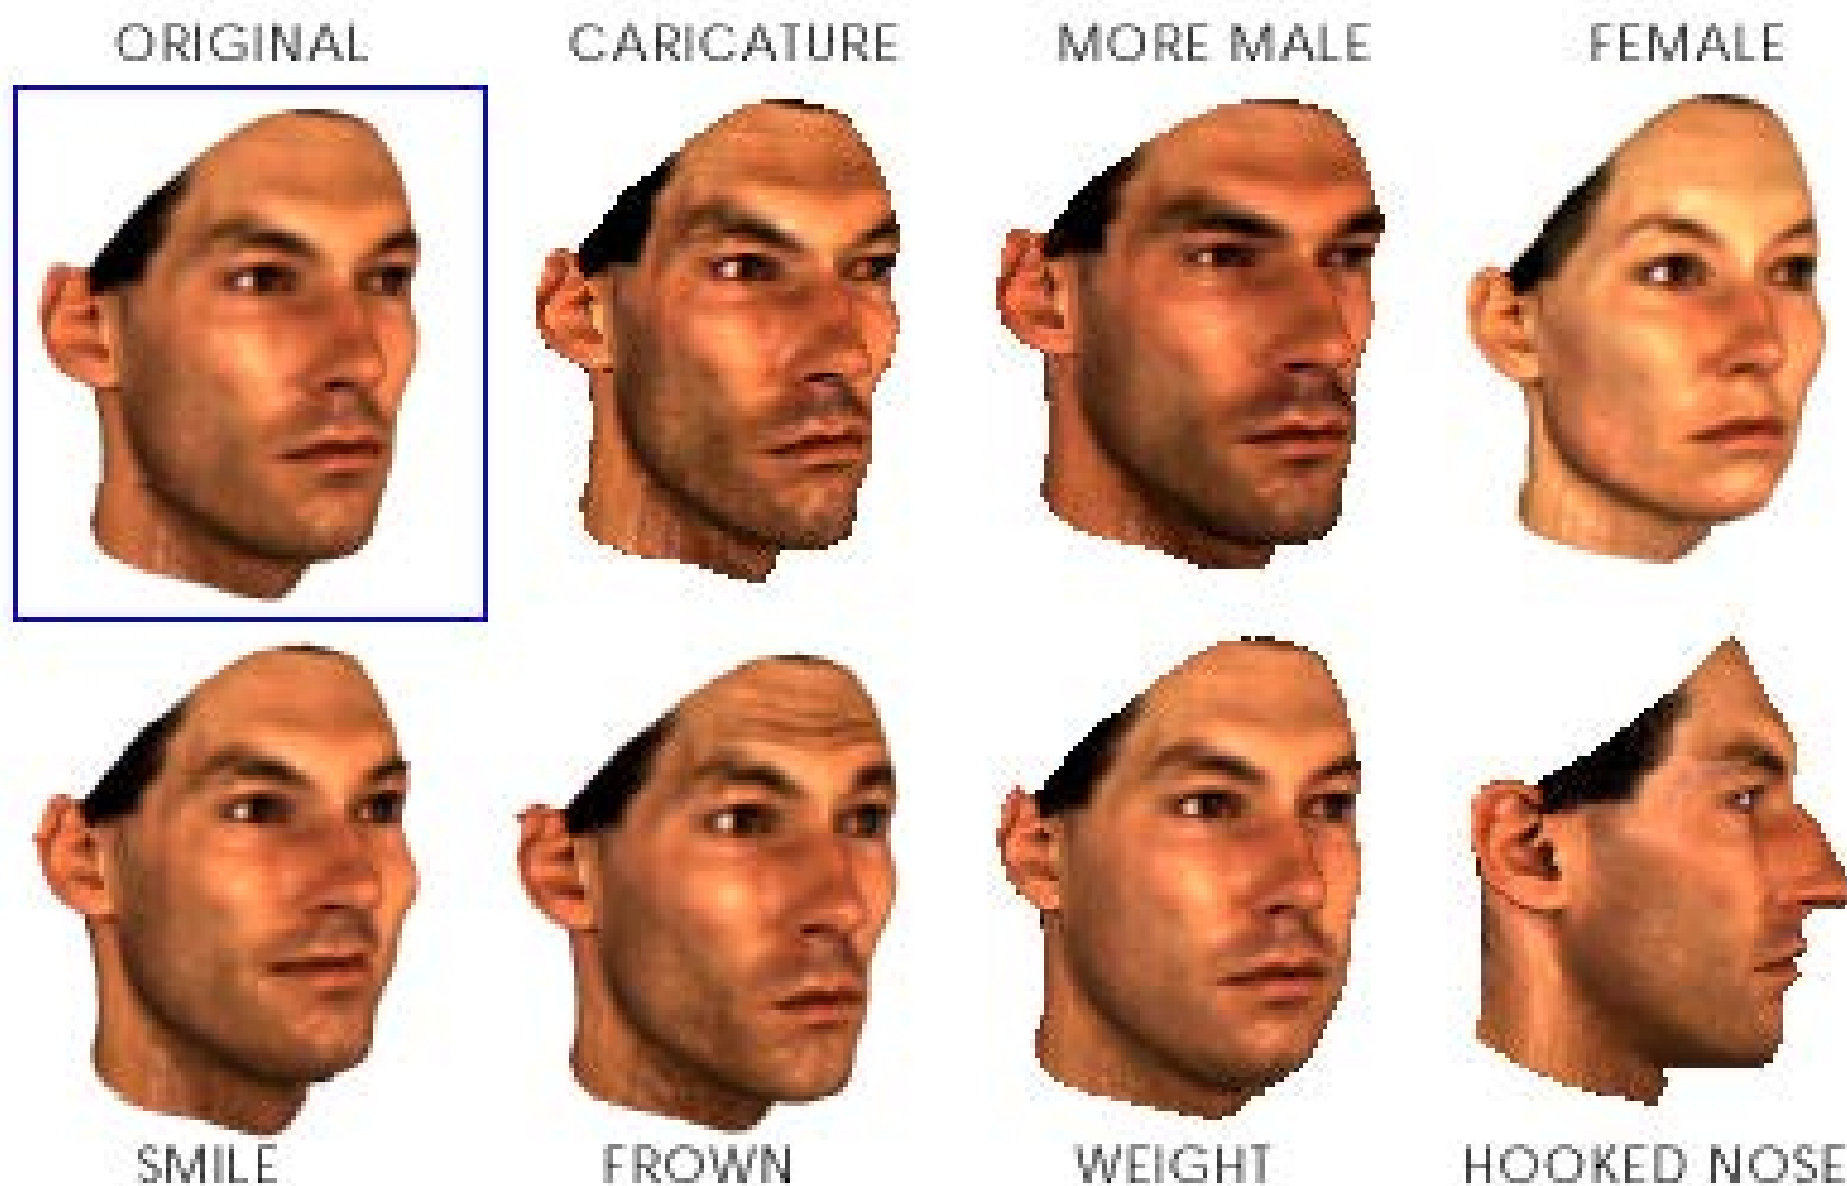
---



- How many eigenfaces to use?
- Look at the decay of the eigenvalues
  - the eigenvalue tells you the amount of variance “in the direction” of that eigenface
  - ignore eigenfaces with low variance

# PCA with depth data: Blinz & Vetter, 1999

---



<http://www.youtube.com/watch?v=jrutZaYoQJo>

# Non-linear Dimensionality Reduction

---

# Learn an embedding (“self-supervision”)

[Collobert and Weston 2008]

True Snippet

NN

house, where the professor lived without his wife and child; or so he said jokingly sometimes: “Here’s where I live. My house.” His daughter often added, without resentment, for the visitor’s information, “It started out to be for me, but it’s really his.” And she might reach in to bring forth an inch-high table lamp with fluted shade, or a blue dish the size of her little fingernail, marked “Kitty” and half full of eternal milk; but she was sure to replace these, after they had been admired, pretty near exactly where they had been. The little house was very orderly, and just big enough for all it contained, though to some tastes the bric-à-brac in the parlor might seem excessive. The daughter’s preference was for the store-bought gimmicks and a little es, the toasters and carpet sweepers of Lilliput, but she knew that most adult visitors would

# Skip Gram (word2vec)


[\[Mikolov et al. 2013\]](#)

house, where the professor lived without his wife and child; or so he said jokingly sometimes: "Here's where I live. My house." His daughter often added, without resentment, for the visitor's information, "It started out to be for me, but it's really his." And she might reach in to bring forth an inch-high table lamp with fluted shade, or a blue dish the size of her little fingernail, marked "Kitty" and half full of eternal milk; but she was sure to replace these, after they had been admired, pretty near exactly where they had been. The little house was very orderly, and just big enough for all it contained, though to some tastes the bric-à-brac in the parlor might seem excessive. The daughter's preference was for the store-bought gimmicks and appliances, the toasters and carpet sweepers of Lilliput, but she knew that most adult visitors would




NN

# Learning word2vec embeddings





Word	Cosine distance
norway	0.760124
denmark	0.715460
finland	0.620022
switzerland	0.588132
belgium	0.585835
netherlands	0.574631
iceland	0.562368
estonia	0.547621
slovenia	0.531408





## Example

$$\text{vec}(\text{"man"}) - \text{vec}(\text{"king"}) + \text{vec}(\text{"woman"}) = \text{vec}(\text{"queen"})$$



# Visual Context Task

